

Operating Systems 2023/2024

TP Class 02 – Processes

Vasco Pereira (vasco@dei.uc.pt)

Dep. Eng. Informática da Faculdade de Ciências e Tecnologia da Universidade de Coimbra

Some slides based on previous versions from Bruno Cabral, Paulo Marques and Luis Silva.

operating system

noun

the collection of software that directs a computer's operations, controlling and scheduling the execution of other programs, and managing storage, input/output, and communication resources.

Abbreviation: OS

Source: Dictionary.com

1 2



9 0

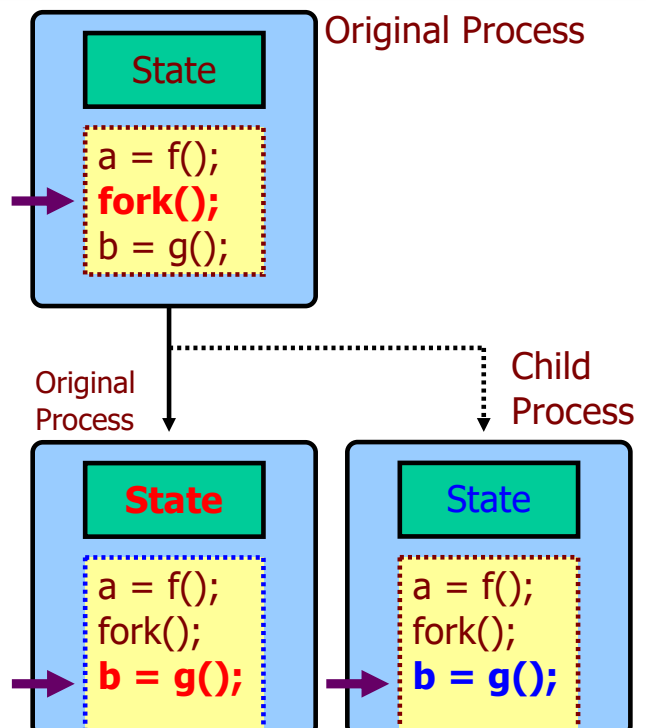
Departamento de
Engenharia Informática
FACULDADE DE
CIÊNCIAS E TECNOLOGIA
UNIVERSIDADE DE
COIMBRA

Updated on
04 February 2024

Operating Systems - University of Coimbra
2

Process Model

- Process creation in Unix is based on spawning child processes which inherit all the characteristics of their fathers
 - Variables, program counter, open files, etc.
 - Spawning a process is done using the `fork()` system call
- After forking, each process will be executing having different variables and different state.
 - The Program Counter will be pointing to the next instruction
 - Changing a variable in the child program doesn't affect its father (and vice-versa)



Process Management

- Each process has an unique identifier (**PID**). Each process has a father, which is also identified (PPID).
- `pid_t getpid(void);`
 - Returns the PID of the current process.
- `pid_t getppid(void);`
 - Returns the PID of the parent process.
- `pid_t fork(void);`
 - Creates a new process which inherits all its father's state. It returns **0 to the child process** and **the child's PID to the original process**.
- `pid_t wait(int* status);`
 - Waits until a child process exits. The status of the child is set in status. (status is the return value of the process). Returns -1 in case of error.
- `pid_t waitpid(pid_t who, int* status, int options);`
 - Same as wait() but allows to wait for a particular child. By using WNOHANG in options, allows for checking if a child has already exited without blocking. 0 in who means "wait for any child". Returns -1 in case of error.

Using processes to do different things

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>

int main()
{
    pid_t id;

    id = fork();
    if (id == 0)
    {
        printf("[%d] I'm the son!\n", getpid());
        printf("[%d] My parent is: %d\n", getpid(), getppid());
    }
    else
    {
        printf("[%d] I'm the father!\n", getpid());
        wait(NULL);
    }
    return 0;
}
```

The key idea is to create asymmetry between processes

And the result is...

```
pmarques@null:~/IPC$ gcc -Wall hello_world.c -o hello_world
pmarques@null:~/IPC$ ./hello_world
[2190] I'm the father!
[2191] I'm the son!
[2191] My parent is: 2190
pmarques@null:~/IPC$ _
```

But why did we do `wait(NULL)`?

```
...
printf("[%d] I'm the father!\n", getpid());
wait(NULL);
...
```

`wait(NULL)` ensures that the father process waits for any of its children, discarding child exit status.

Process Termination in UNIX

- A process is only truly eliminated by the operating system when its father calls `wait()/waitpid()` on it.
 - This allows the parent check things like the exit code of its son's
- **Zombie Process:** One that has died and its parent has not acknowledged its death (by calling `wait()`)
 - Be careful with this if your are designing servers. They are eating up resources!!
- **Orphan Process:** One whose original parent has died. In that case, its parent becomes *init* (process 1). (not always true; in some recent versions of operating systems the parent becomes the nearest sub-reaper process - a sub-reaper fulfills the role of *init*(1) for its descendant processes.)

Let's generate some Zombies

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>

void worker() {
    printf("[%d] Hi, I'm a worker process! Going to die...\n",
        getpid());
}

int main()
{
    for (int i=0; i<10; i++) {
        if (fork() == 0) {
            worker();
            exit(0);
        }
    }
    printf("[%d] Big father is sleeping!\n", getpid());
    sleep(10);

    return 0;
}
```

demo2.c

Let's see adoption

```
#include (...)

void worker()
{
    sleep(10);
    printf("[%d] Let's see who my dad is: %d\n", getpid(),
        getppid());
}

int main()
{
    for (int i=0; i<10; i++)
    {
        if (fork() == 0)
        {
            worker();
            exit(0);
        }
    }

    printf("[%d] Big dad is going away!\n", getpid());
    return 0;
}
```

demo3.c

How to structure code


```
(...)  
if ((id = fork()) == 0)  
{  
    // Huge amount of code  
    // ...  
}  
else  
{  
    // Huge amount of code  
    // ...  
}  
(...)
```

Don't do it!

Fairly common...

How to structure code

```
void client(int id)  
{  
    // Client code  
    // ...  
}  
  
if ((id = fork()) == 0)  
{  
    client(id);  
    exit(0);  
}  
else if (id == -1)  
{  
    error();  
}  
  
// Original process code  
// ...
```



Note: You still have to consider how to take care of zombies

How a process becomes another executable


- Somehow the OS must be able to execute code starting from an executable file
 - e.g. how does the shell (bash) becomes 'ls'?
- `int execl(const char *path, const char *arg, ...);`
- `int execlp(const char *file, const char *arg, ...);`
- `int execlxe(const char *path, const char *arg, ..., char *const envp[]);`
- `int execv(const char *path, char *const argv[]);`
- `int execvp(const char *file, char *const argv[]);`
- “exec family” of functions
 - Allow to substitute the current process executable image by another one. The substitution is complete!
 - The functions that have a ‘p’ make use of the environment PATH; The functions that have a ‘v’ make use of a pointer to an array on parameters; The functions that have an ‘l’ have the parameters passed separated by commas
 - Make sure that the first parameter is the name of the program!

Example

- Simple program that lists the files in the current directory

```
pmarques@null:~/IPC$ g++ -Wall list_files.c -o list_files
pmarques@null:~/IPC$ ./list_files
.      adopt.c      list_files      posix_version      zombie.c
..     hello_world  list_files.c   posix_version.c
adopt  hello_world.c  posix         zombie
pmarques@null:~/IPC$ _
```

- Note: A successful `exec()` never returns
 - The code, the stack, the heap, it's all replaced by the new executable



The corresponding code...

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>

int main()
{
    if (execlp("ls", "ls", "-a", NULL) == -1)
        perror("Error executing ls: ");
    else
        printf("This cannot happen!\n");

    return 0;
}
```

demo4.c

Using an array
can be more
flexible...

```
char* ls_param[] = { "ls", "-a", NULL };

if (execvp(ls_param[0], ls_param) == -1)
    perror("Error executing ls: ");
```

Code to wait for the end of an exec call...

```
makeargv(buf, DELIM, &args);
pid_t child = fork();
if (child == 0)
{
    execvp(args[0], args);
    perror("Erro ao executar o comando");
    exit(-1);
}
waitpid(child, NULL, 0);
```


Information about processes in Linux

- **Command `ps` (process status)**
 - reports a snapshot of the current processes
 - To see all processes in the OS in full-format listing
`ps -ef`
 - To select all processes from user “user1”
`ps -ef | grep user1`
 - **Some process state codes shown in `ps`**
 - Seen when `ps` is used with a state output modifier
 - E.g.: select all processes and show specific columns
 - `ps -eo uid,pid,ttty,time,cmd,stat`
 - **R** running
 - **T** stopped by a job control signal or because it is being traced
 - **Z** defunct (“zombie”)
 - **+** the process is in the foreground process group.

Class demos included


- **Demo01 – create process**
`demo01.c`
- **Demo02 – zombies**
`demo02.c`
- **Demo02 – adoption**
`demo03.c`
- **Demo04 – exec**
`demo04.c`

References



- **Advanced Programming in the UNIX Environment**
3rd Edition (2013)
W. Richard Stevens, Stephen A. Rago
Addison-Wesley
 - Chapter 8: Process Control

Where to learn more?



- **The Linux Programming Interface**
2010
Michael Kerrisk
No Starch Press
 - Chapters 24 to 27

INTRODUCTION TO ASSIGNMENT 03 – “PROCESSES”

Thank you! Questions?



I keep six honest serving men. They taught me all I knew. Their names are What and Why and When and How and Where and Who.
—Rudyard Kipling