

# Redes de Comunicação 2023/2024

---

## T04 Network layer

Jorge Granjal  
University of Coimbra



# T04: outline

## 4.1 Overview of Network layer

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- service model
- datagram format
- fragmentation
- IPv4 addressing
- network address translation


## 4.4 Routing protocols

- link state
- Distance vector
- Dijkstra and RIP

## 4.5 ICMP: The Internet Control Message Protocol

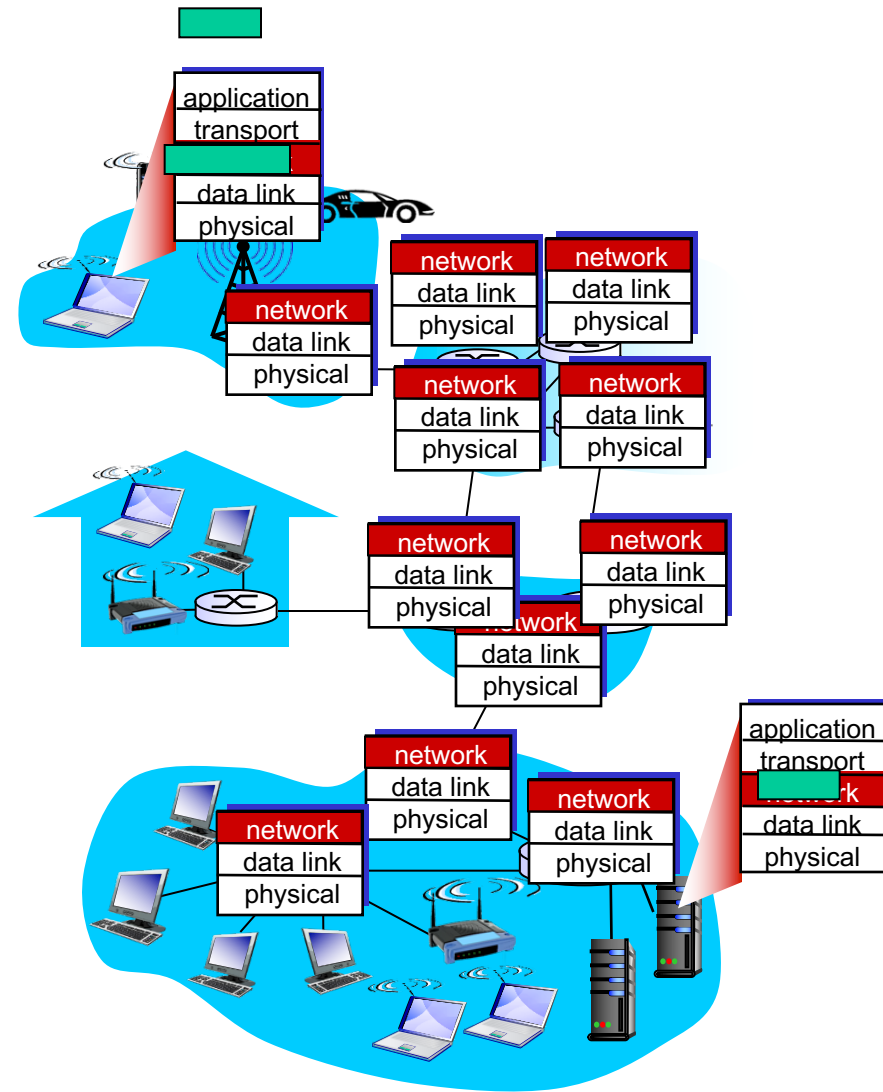
# Internet (TCP/IP) protocol stack

- *application*: supporting network applications
  - FTP, SMTP, HTTP, ...
- *transport*: process-process data transfer
  - TCP, UDP
- *network*: routing of datagrams from source to destination
  - IP, routing protocols
- *link*: data transfer between neighboring network elements
  - Ethernet, 802.111 (WiFi), PPP
- *physical*: bits “on the wire”



# Network layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it



# Two key network-layer functions

## *network-layer functions:*

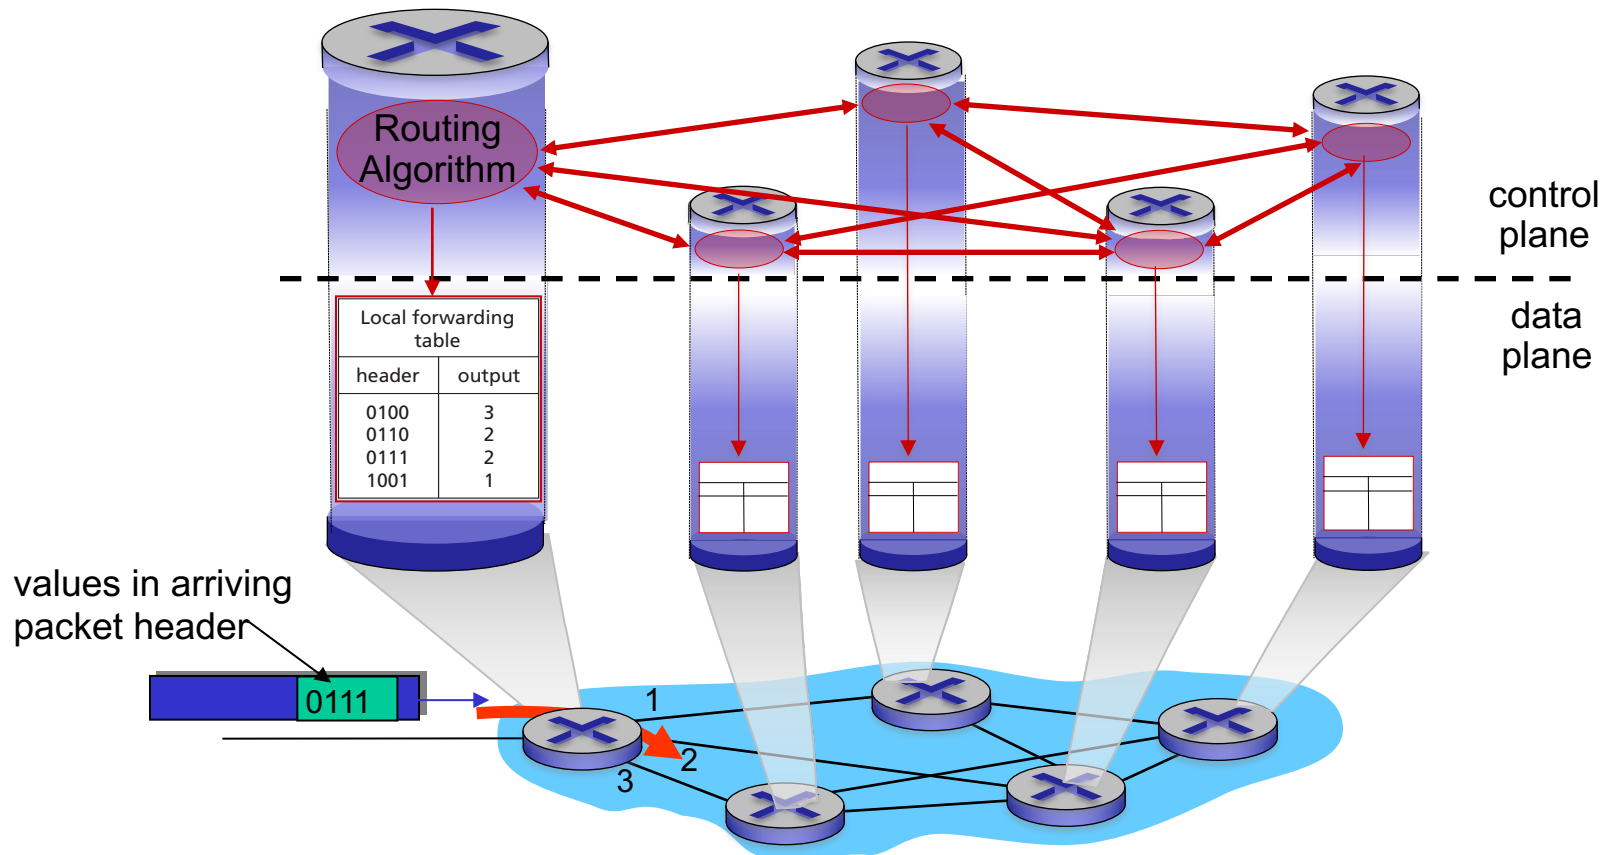
- *forwarding*: move packets from router's input interface to appropriate router output interface
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*
  - *implemented in various routers*

## *analogy: taking a trip*

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination

# Routing algorithms and forwarding tables

Individual routing algorithm components *in each and every router* interact in the control plane



# Routing algorithms: important aspects

- Algorithms compute the *forwarding tables* that are used to forward packets through the network
- Router receives *routing control messages* which are used to configure its *forwarding table*
- Routing algorithms may be:
  - Centralized: algorithms executes on central site and routing information is downloaded to each of the routers
  - Decentralized: a piece of the routing algorithm runs in each router
- The term “*packet switch*” may refer to either:
  - Link-layer switch: forwarding decision is based on fields of the link-layer frame (layer 2 devices)
  - Router or layer 3 switch: forwarding decision is based on fields of the network-layer (IP) header

# T04: outline

## 4.1 Overview of Network layer

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- service model
- datagram format
- fragmentation
- IPv4 addressing
- network address translation

## 4.4 Routing protocols

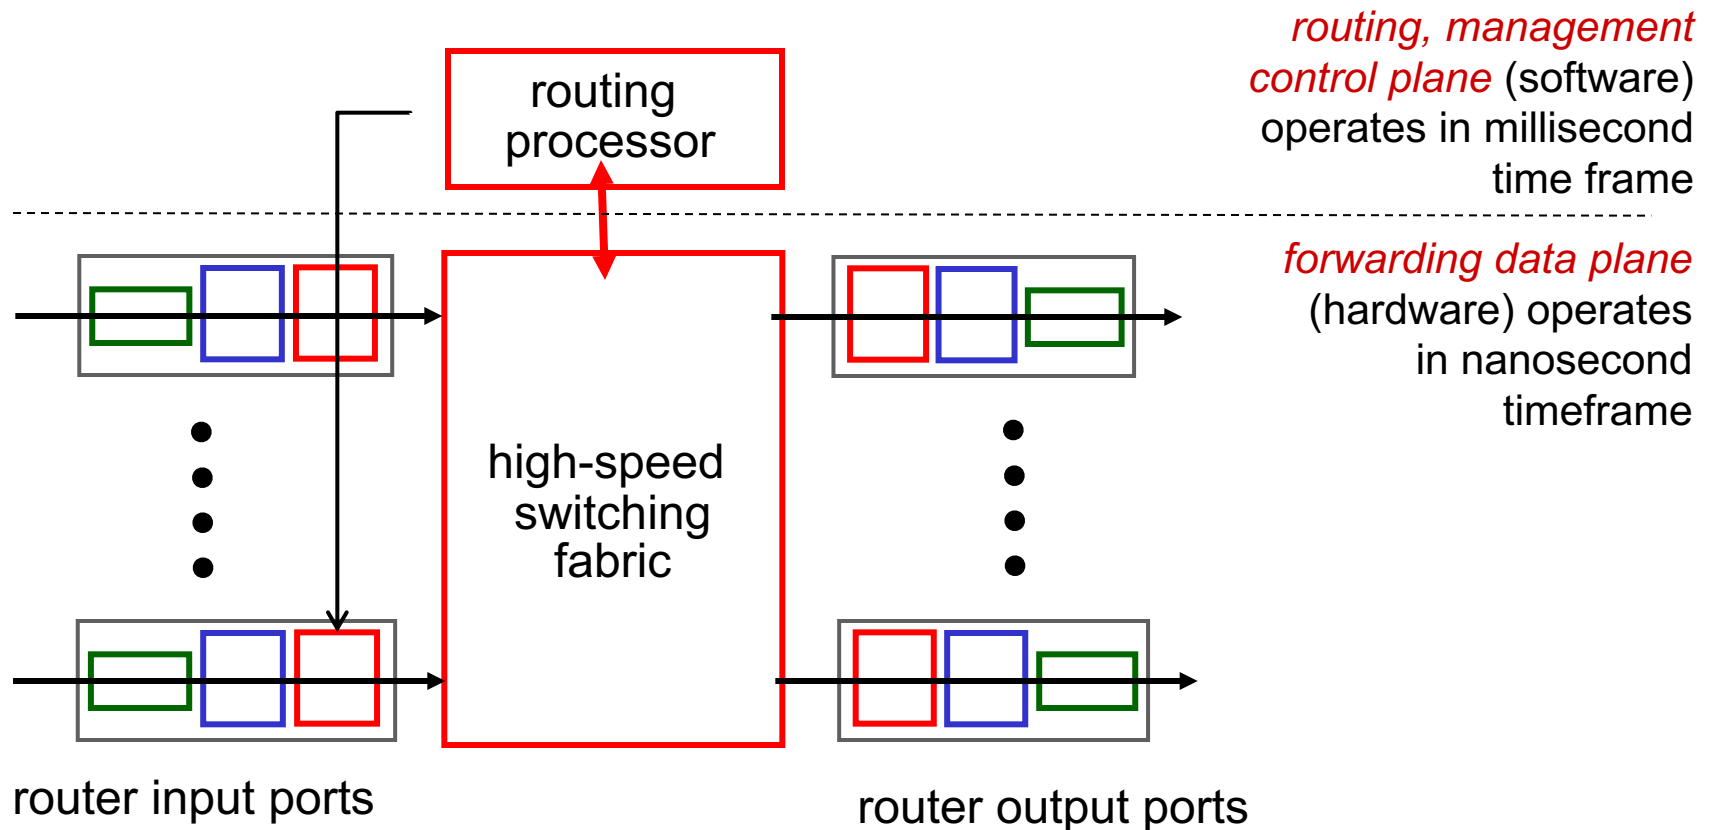
- link state
- Distance vector
- Dijkstra and RIP

## 4.5 ICMP: The Internet Control Message Protocol



# Router architecture overview

- high-level view of generic router architecture:



# Destination-based forwarding

*forwarding table*

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

# Longest prefix matching

## *longest prefix matching*

when looking for forwarding table entry for given destination address, use *longest* address prefix (longest subnet mask) that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

DA: 11001000 00010111 00010110 10100001


which interface?

DA: 11001000 00010111 00011000 10101010

which interface?

# Scheduling mechanisms

- *scheduling*: choose next packet to send on link
- *FIFO (first in first out) scheduling*: send in order of arrival to queue
  - *discard policy*: if packet arrives to full queue: who to discard?
    - *tail drop*: drop arriving packet
    - *priority*: drop/remove on priority basis
    - *random*: drop/remove randomly



# T04: outline

4.1 Overview of Network layer

4.2 What's inside a router

4.3 IP: Internet Protocol

- service model
- datagram format
- fragmentation
- IPv4 addressing
- network address translation

4.4 Routing protocols

- link state
- Distance vector
- Dijkstra and RIP

4.5 ICMP: The Internet Control Message Protocol

# Internet's network layer service

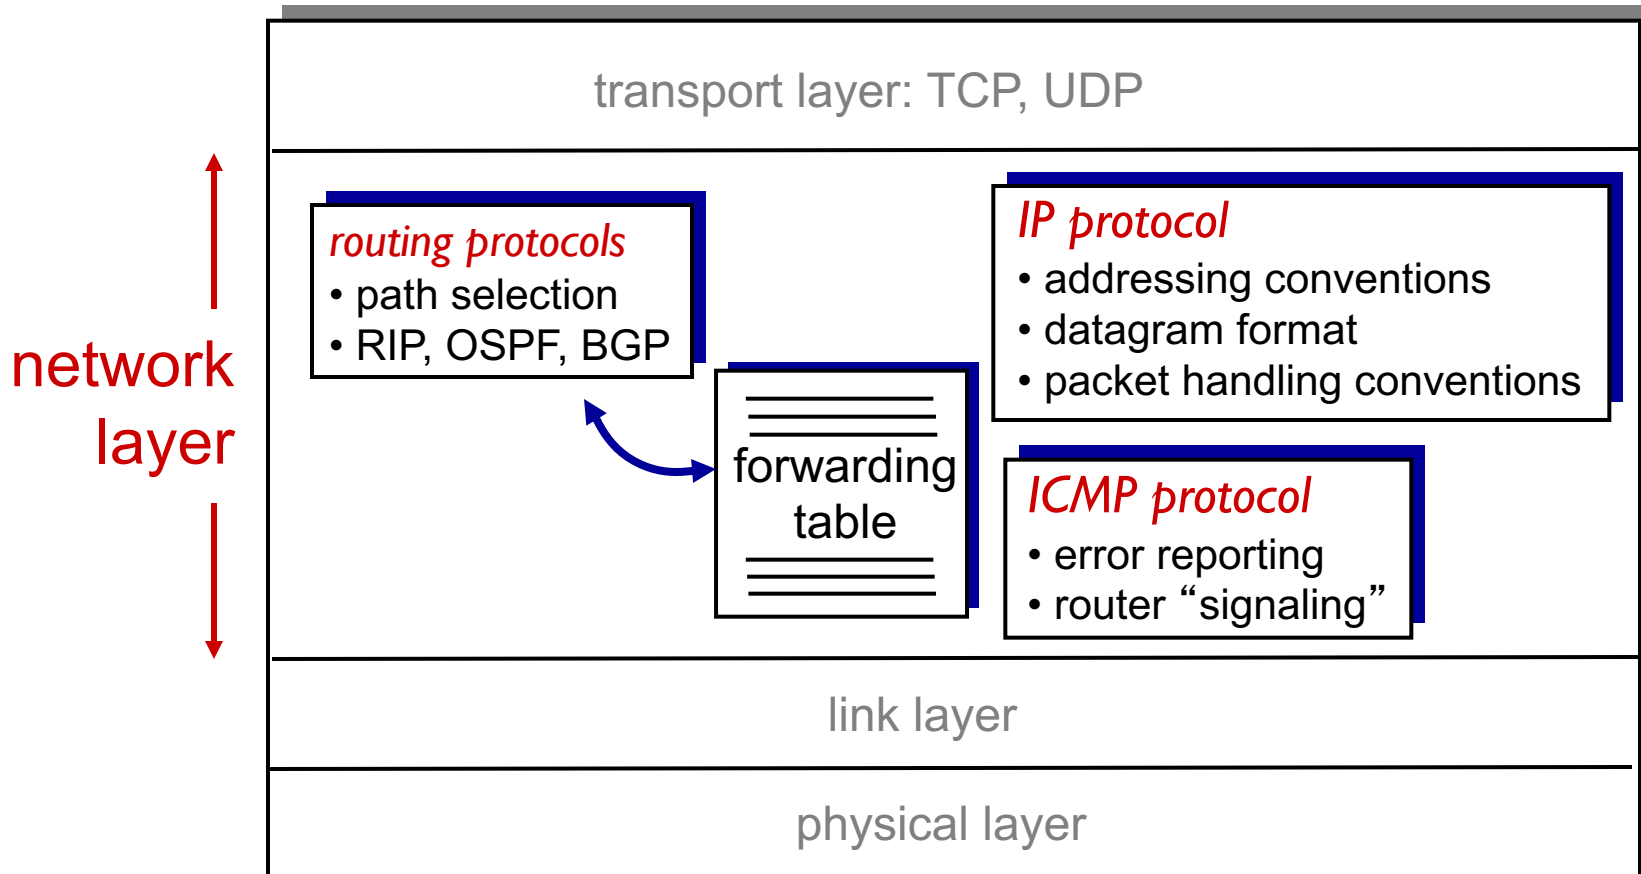
Internet's network layer offers a **best-effort** service:



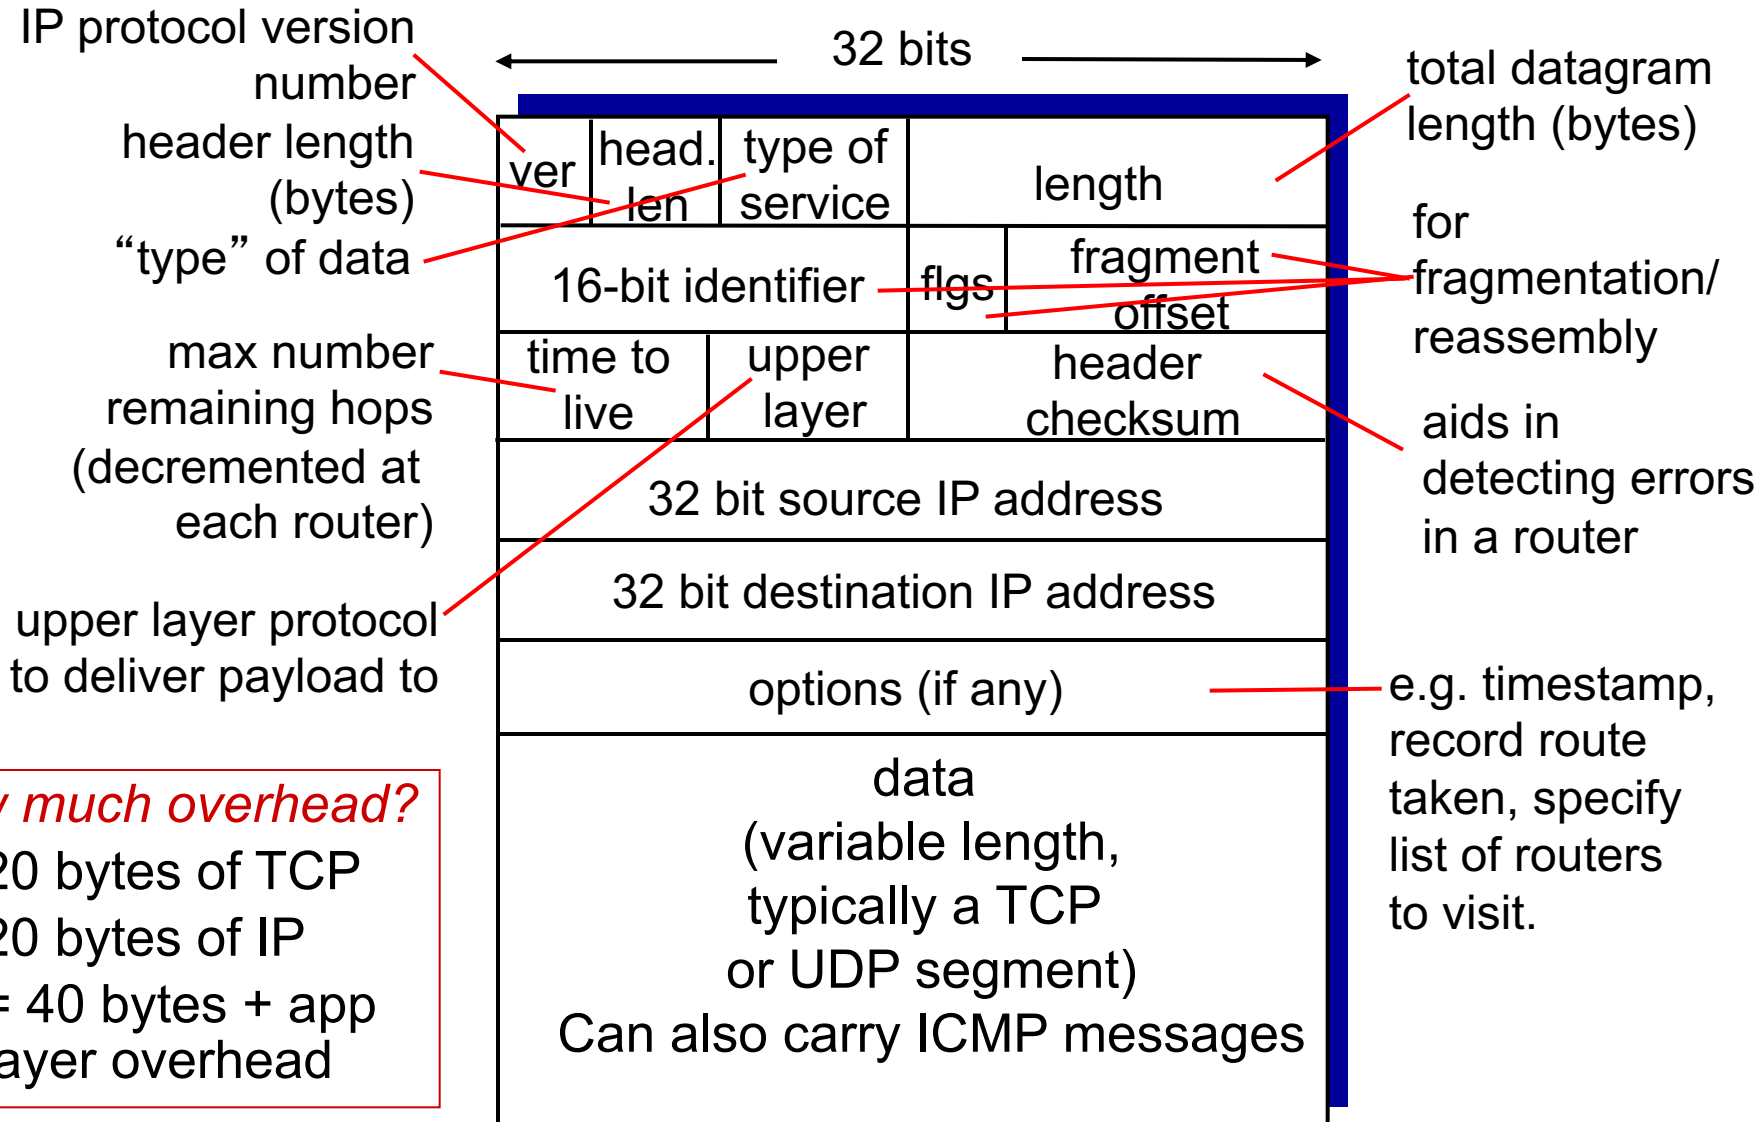
- Service model: best effort
- Bandwidth guarantee: None
- No-Loss guarantee: None
- Ordering: any order possible
- Timing between packets: not maintained
- Congestion indication: None

# The Internet network layer

host, router network layer functions:



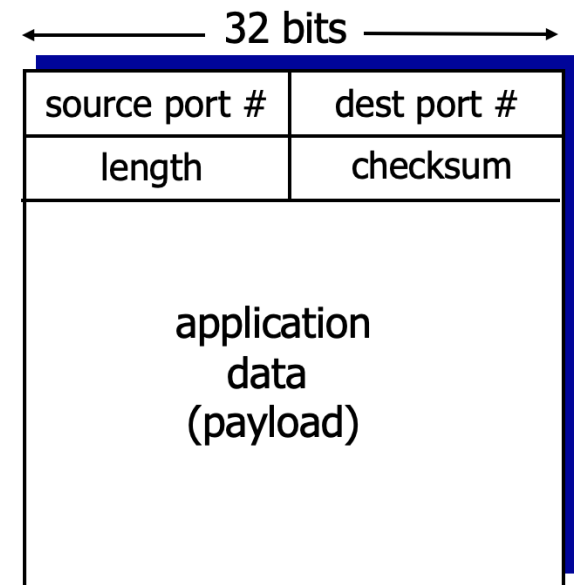
# IP datagram format





# Revisiting UDP: packet size

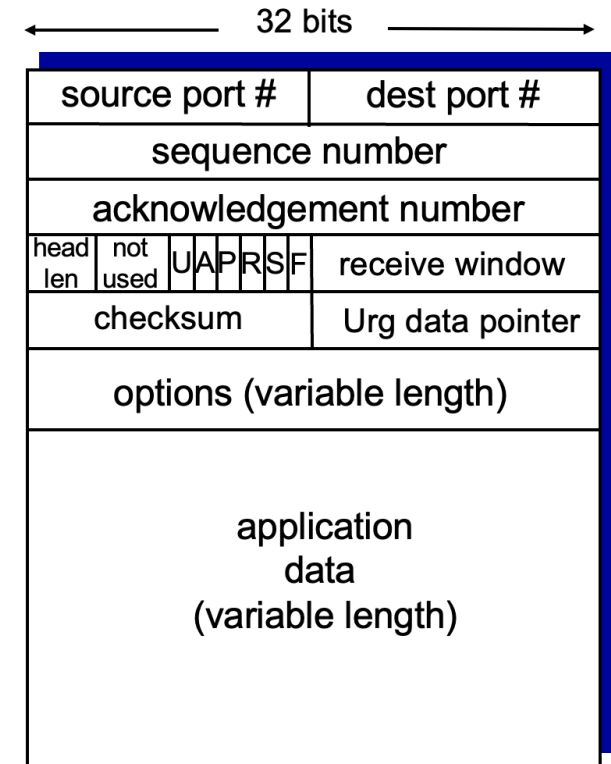
- Minimum size of a UDP packet is 8 bytes.
- The maximum size of a UDP packet is 65,535 bytes (or  $2^{16} - 1$  bytes) minus the size of the IP (20 bytes) and UDP (8 bytes) headers.
- Most networks have a lower maximum transmission unit (MTU) size, which limits the maximum size of a UDP packet that can be transmitted without fragmentation.
- It is generally recommended to keep UDP packets small to minimize the chances of packet loss.



UDP segment format

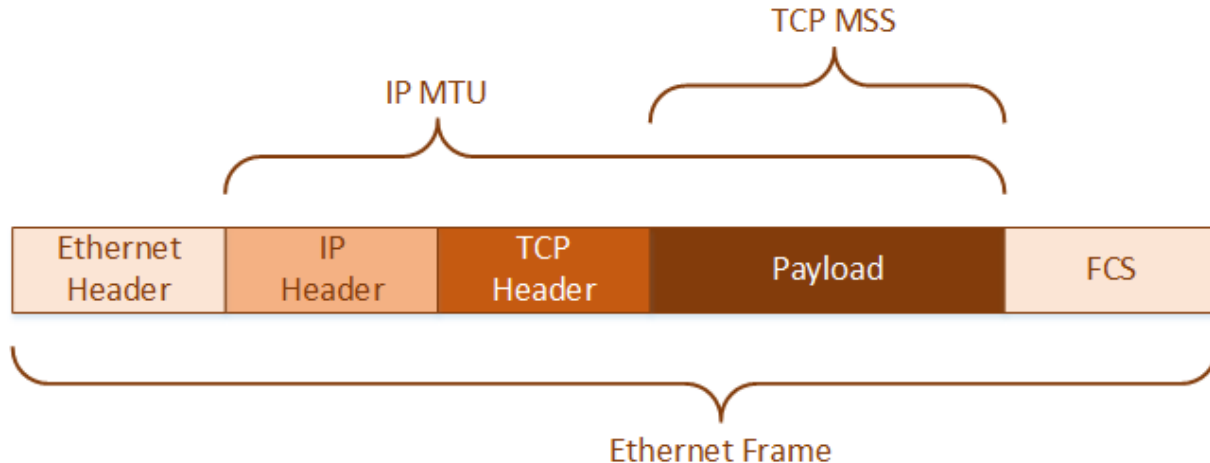
# Revisiting TCP: packet size

- The minimum size of a TCP packet is 20 bytes.
- The maximum size of a TCP packet is determined by the Maximum Segment Size (MSS) value negotiated between the two endpoints during the TCP connection setup process.
- The MSS value is typically determined by the maximum size of the data that can be transmitted without fragmentation, based on the MTU of the underlying network.
- In practice, the maximum size of a TCP packet can range from 536 bytes (the minimum MSS value) to around 65,535 bytes (the maximum TCP payload size), depending on the MTU of the network and the negotiated MSS value.



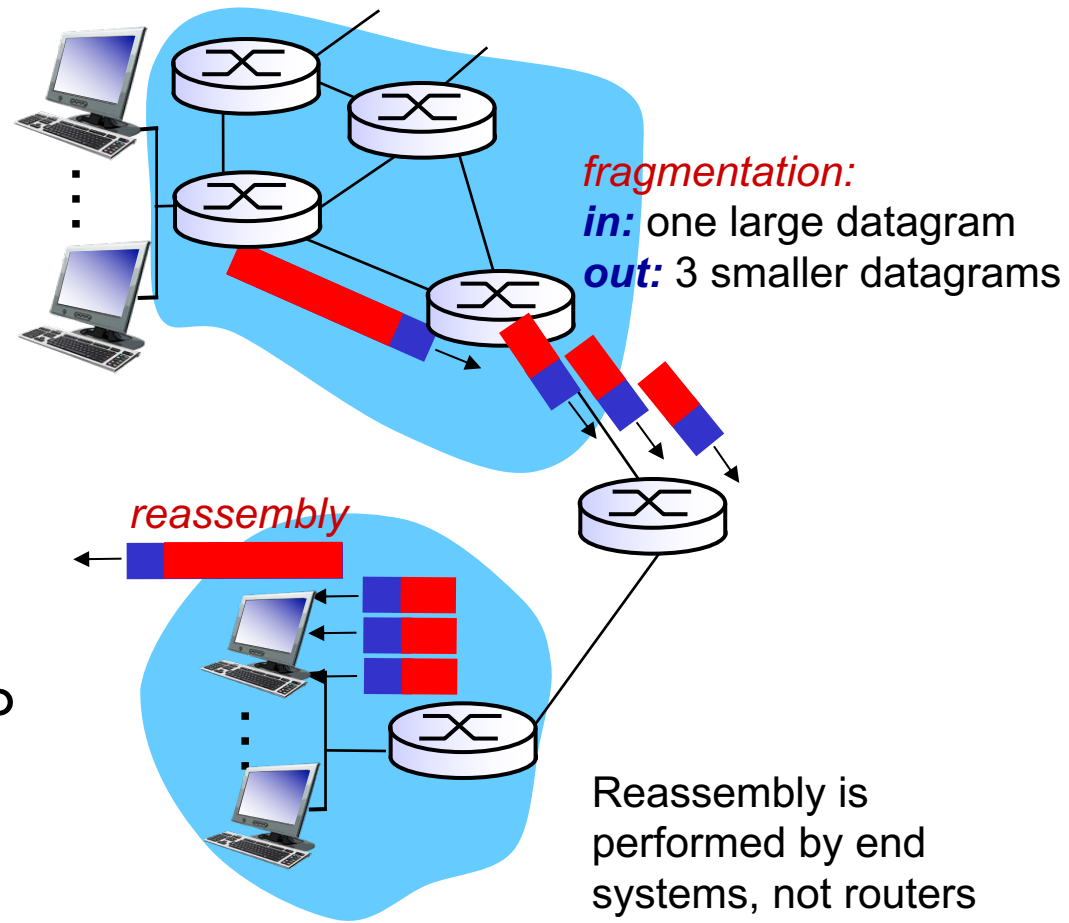
# MTU in TCP

- The maximum size of an IP packet is determined by the Maximum Transmission Unit (MTU) of the underlying network.
- The MTU is the largest size of a packet that can be transmitted over a network without being fragmented. In practice, the MTU can vary depending on the type of network and the network equipment being used.
- For example, the MTU for Ethernet networks is typically 1,500 bytes, which means that the maximum size of an IP packet (including header) that can be transmitted over an Ethernet network is 1,500 bytes.



# IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at final destination (UDP and TCP are expecting to received complete segments)
  - IP header bits used to identify, order related fragments



# IP fragmentation, reassembly

## *example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

*one large datagram becomes  
several smaller datagrams*

1480 bytes in  
data field

offset =  
 $1480/8$

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

# IP addresses: how to get one?

**Q:** How does a *host* get IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
  - “plug-and-play”

# DHCP: Dynamic Host Configuration Protocol


*goal:* allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

## *DHCP overview:*

- host broadcasts “DHCP discover” msg [optional]
- DHCP server responds with “DHCP offer” msg [optional]
- host requests IP address: “DHCP request” msg
- DHCP server sends address: “DHCP ack” msg

# DHCP client-server scenario






# DHCP client-server scenario

DHCP server: 223.1.2.5

DHCP discover

arriving  
client



Broadcast: is there a  
DHCP server out there?

DHCP offer

Broadcast: I'm a DHCP  
server! Here's an IP  
address you can use

DHCP request

Broadcast: OK. I'll take  
that IP address!

DHCP ACK


Broadcast: OK. You've  
got that IP address!

# DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:


- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

# DHCP: example



- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

# DHCP: example



- DCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

# IP addresses: how to get one?

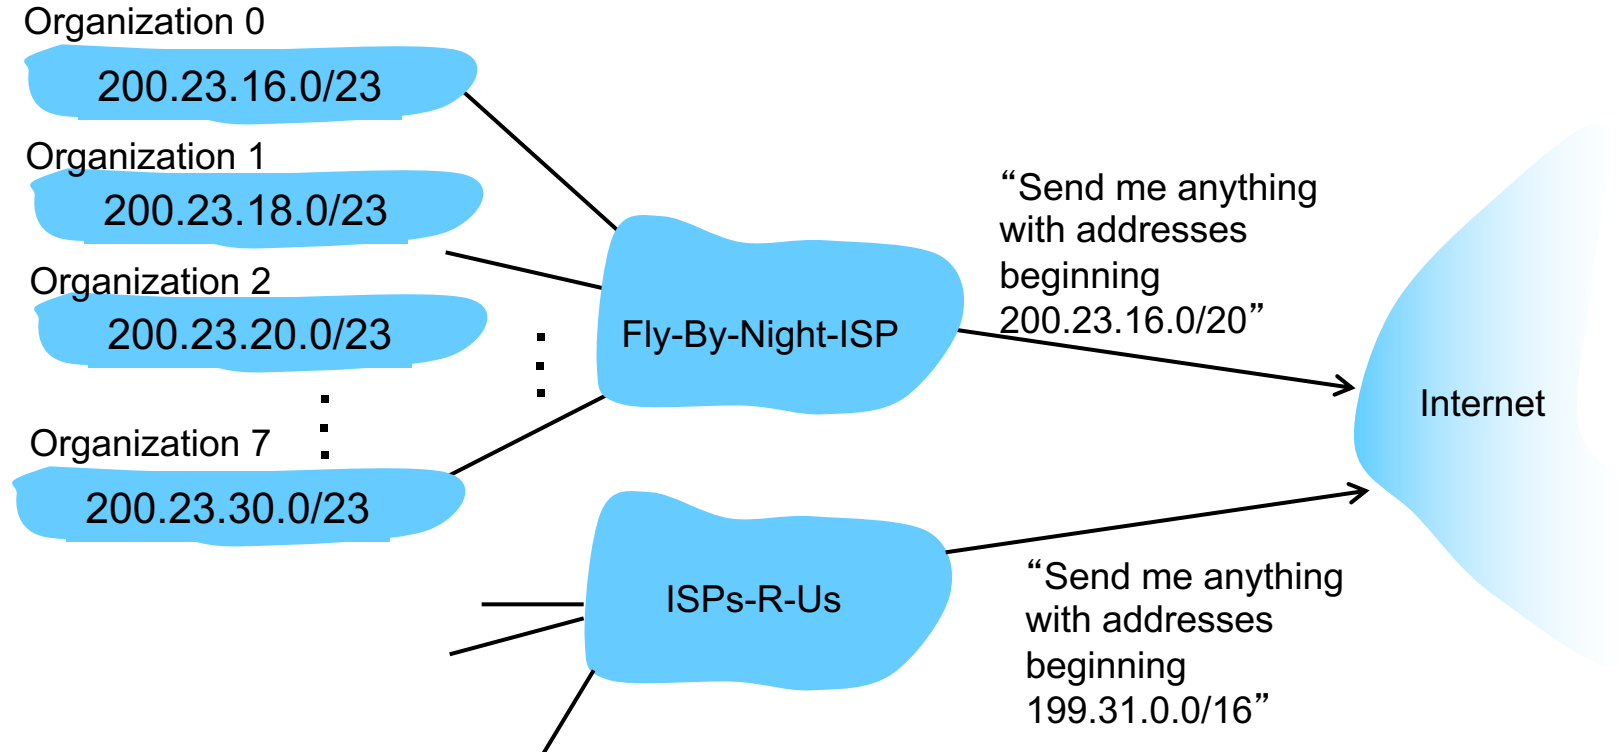
**Q:** how does *network* get subnet part of IP addr?

**A:** gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>00010010</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>00010100</u>	00000000	200.23.20.0/23
...	.....			....	....
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>00011110</u>	00000000	200.23.30.0/23

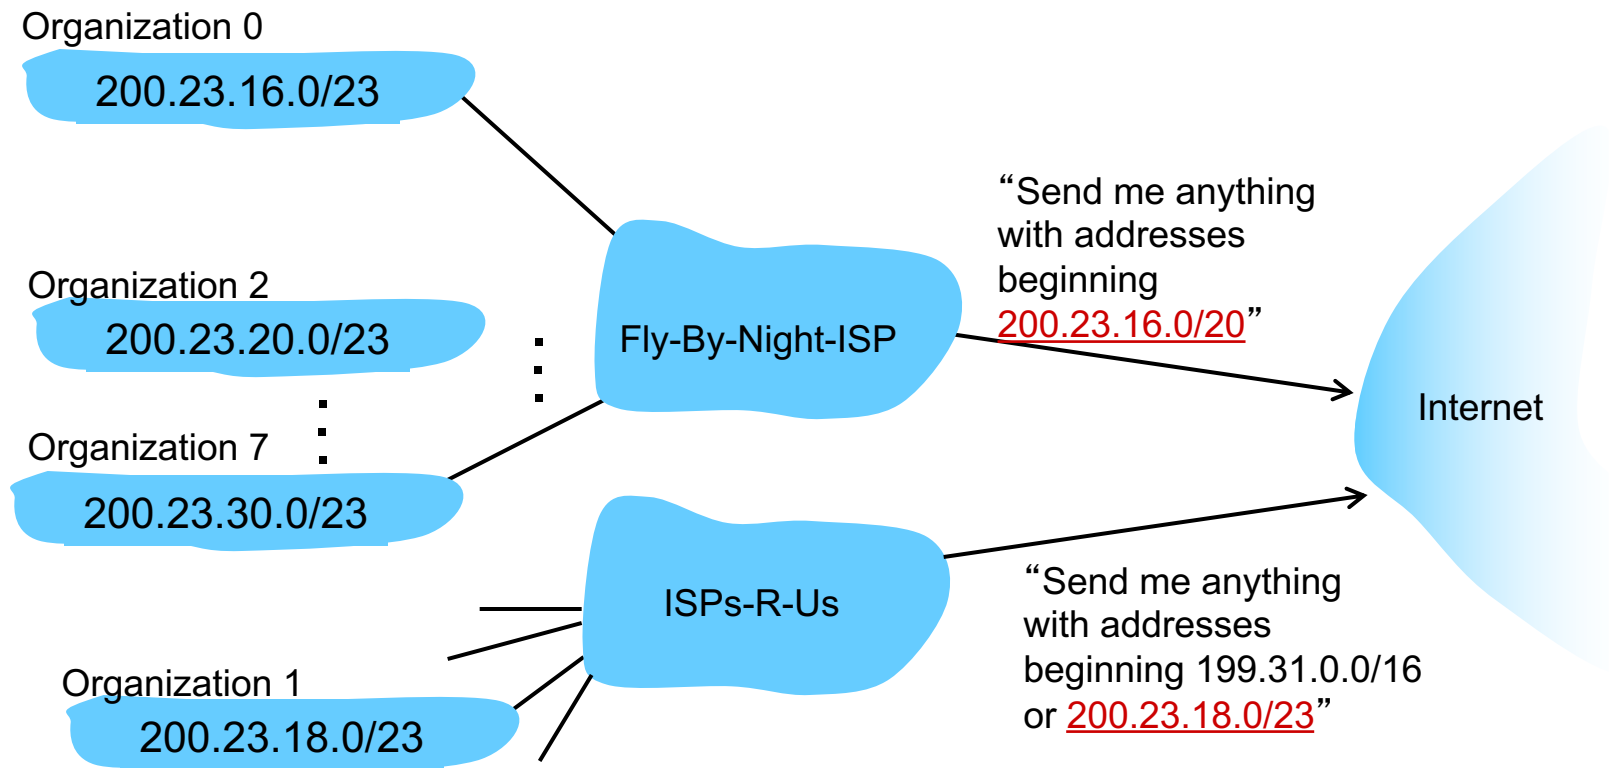
# Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:



# Hierarchical addressing: more specific routes

ISPs-R-Us has a more specific route to Organization 1



# IP addressing: the last word...

**Q:** how does an ISP get block of addresses?

**A: ICANN:** Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes





# NAT: private IP address ranges

Private address range		
Class	start address	finish address
A	10.0.0.0	10.255.255.255
B	172.16.0.0	172.31.255.255
C	192.168.0.0	192.168.255.255

Public address range		
Class	start address	finish address
A	0.0.0.0	126.255.255.255
B	128.0.0.0	191.255.255.255
C	192.0.0.0	223.255.255.255
D	224.0.0.0	239.255.255.255
E	240.0.0.0	254.255.255.255

# NAT: network address translation

*motivation:* local network uses just one IP address as far as outside world is concerned:

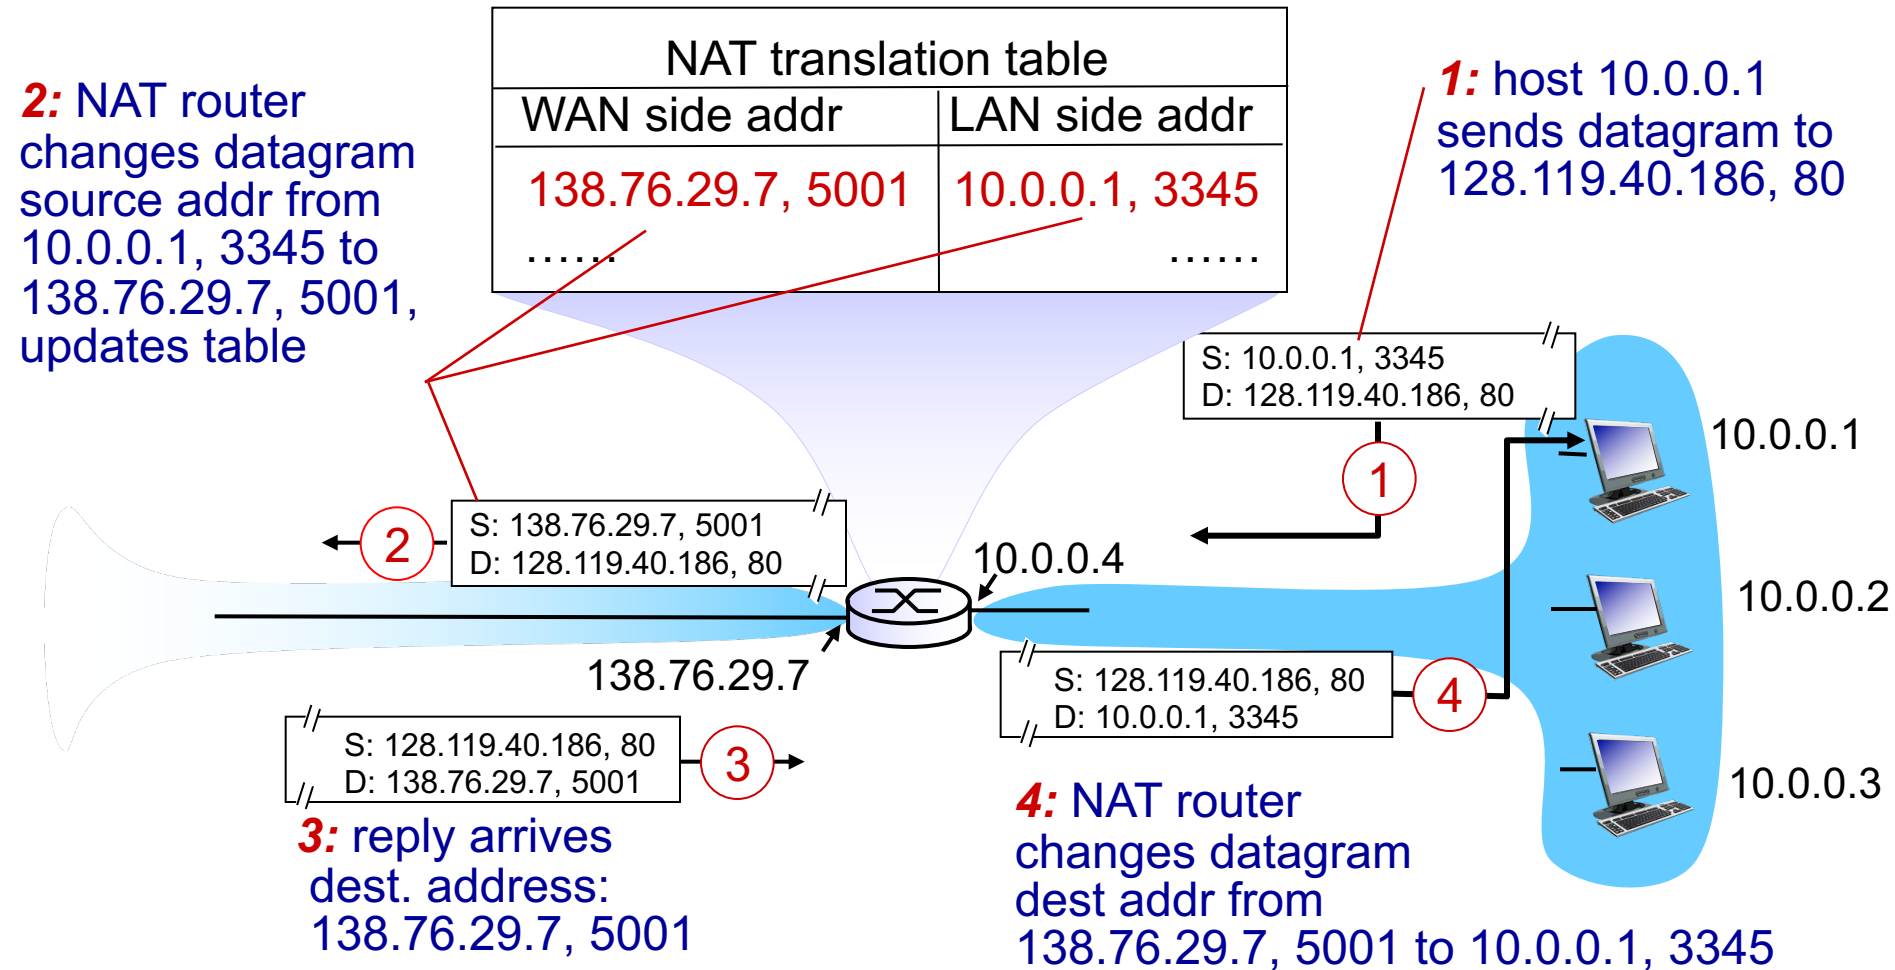
- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

# NAT: network address translation

*implementation:* NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)  
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

# NAT: SNAT



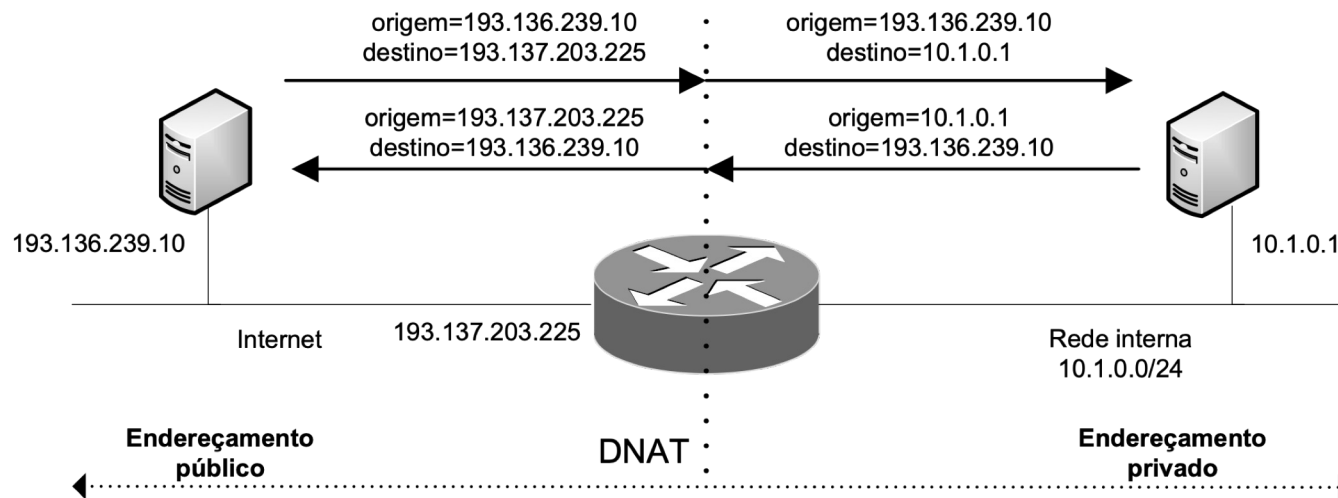
# NAT: SNAT

---

- Also known as SNAT (Source NAT)
- 16-bit port-number field:
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - routers should only process up to layer 3
  - address shortage should be solved by IPv6
  - violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - NAT traversal: what if client wants to connect to server behind NAT?

# NAT Traversal: DNAT

- What if client wants to connect to server behind NAT?
- DNAT (Destination NAT): destination IP and port # changed for packets from the external to the internal network



# T04: outline

## 4.1 Overview of Network layer

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- service model
- datagram format
- fragmentation
- IPv4 addressing
- network address translation

## 4.4 Routing protocols

- link state
- Distance vector
- Dijkstra and RIP

## 4.5 ICMP: The Internet Control Message Protocol


# Network-layer functions

*Recall: two network-layer functions:*

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to destination

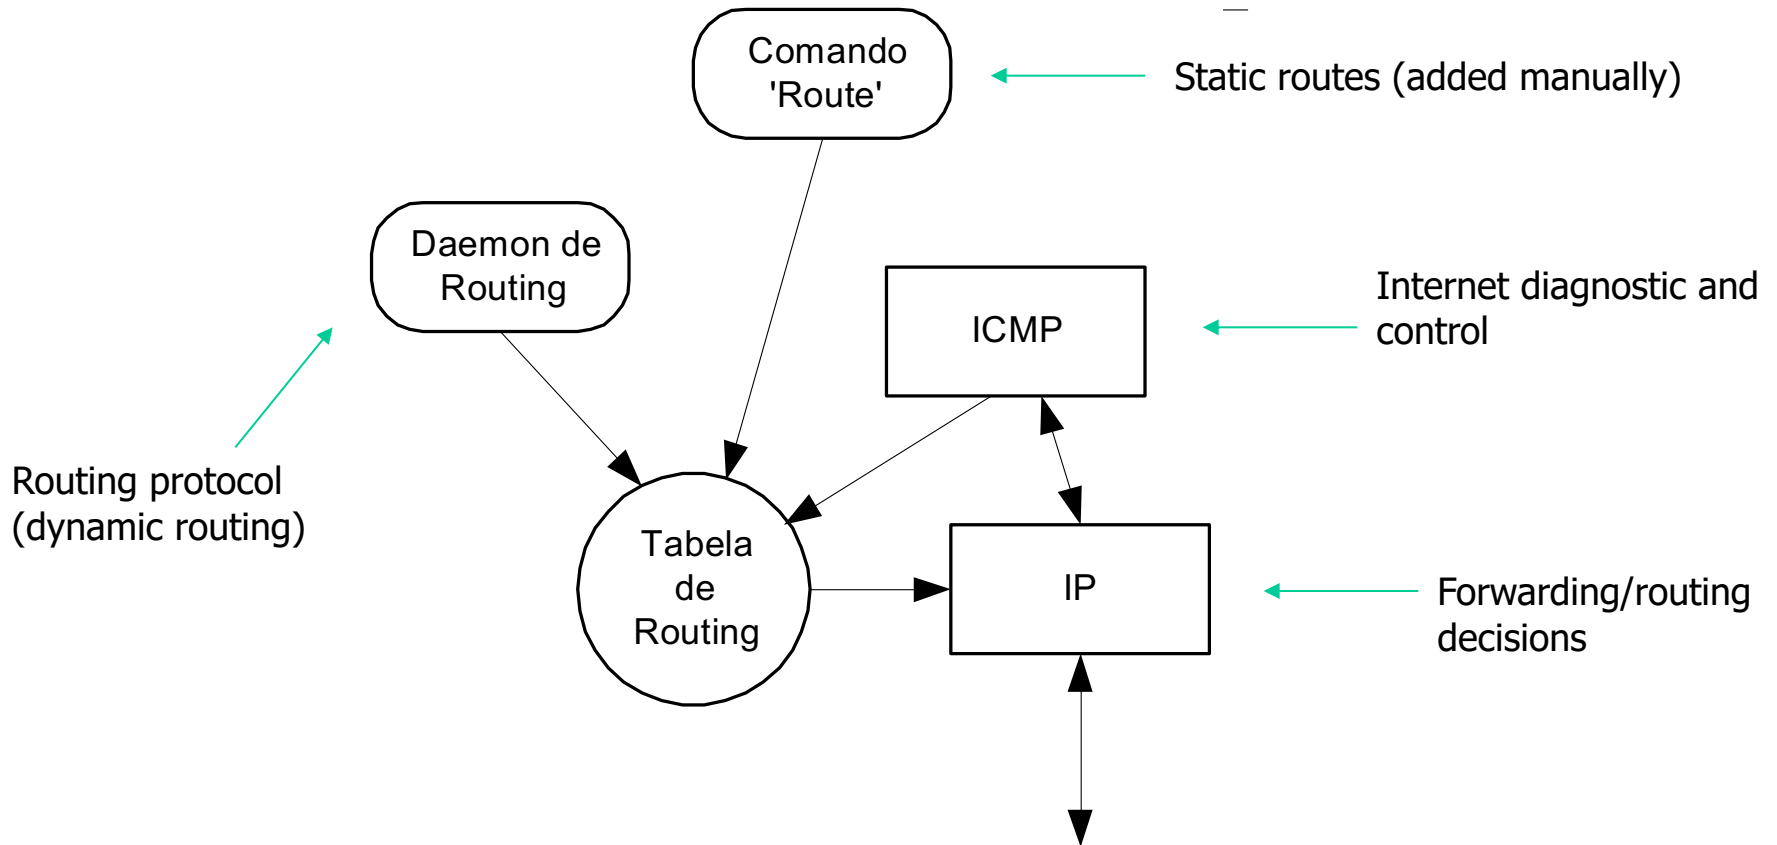
*data plane*

*control plane*





# IP Routing



# Routing table (Cisco)

```
gta#sh ip route
```

```
Codes: C - connected, S - static, I - IGRP, R - RIP, M - mobile
```

```
B - BGP, D - EIGRP, EX - EIGRP external, O - OSPF
```

```
IA - OSPF inter area, N1 - OSPF NSSA external type 1
```

```
N2 - OSPF NSSA external type 2, E1 - OSPF external type 1
```

```
E2 - OSPF external type 2, E - EGP, i - IS-IS
```

```
L1 - IS-IS level-1, L2 - IS-IS level-2, ia - IS-IS inter a
```

```
* - candidate default, U - per-user static route, o - ODR
```


```
P - periodic downloaded static route
```

```
Gateway of last resort is 193.136.094.1 to network 0.0.0.0
```


```
C    193.136.201.0/24 is directly connected, Ethernet0
```

```
R    192.136.200.0/24 [120/1] via 193.136.201.2, 00:00:17, Ethernet0
```


```
R    192.136.202.0/24 [120/1] via 192.138.204.2, 00:00:53, Serial0
```




Origin of routing  
information




Destination address  
and number of bits  
in netmask



Administrative distance  
and metric for route



Gateway (*next hop*)  
to destination



Time since last  
route update



Output  
interface

# Routing algorithm classification

*Q: global or decentralized information?*

*global:*

- all routers have **complete topology**, link cost info
- “link state” algorithms

*decentralized:*

- router knows **only physically-connected neighbors** and link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

*Q: static or dynamic?*


*static:*

- routes change slowly over time

*dynamic:*

- routes change more quickly
  - periodic update
  - in response to link cost or status changes

# Graph abstraction of the network




graph:  $G = (N, E)$

$N$  = set of routers or vertices =  $\{ u, v, w, x, y, z \}$

$E$  = set of links =  $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

*aside:* graph abstraction is useful in other network contexts, e.g., P2P, where  $N$  is set of peers and  $E$  is set of TCP connections

# Graph abstraction: costs



$c(x, x') = \text{cost of link } (x, x')$   
e.g.,  $c(w, z) = 5$


- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion, etc.
- costs can also be different for the two directions of communications

cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

**key question:** what is the least-cost path between u and z ?  
**routing algorithm:** algorithm that finds that least cost path


# Spanning tree (of a graph)

A **spanning tree** is a subset of a graph which has all the vertices covered and with minimum possible number of edges



# Dijkstra algorithm

- Created by Edsger Dijkstra (1959)
- Algorithm to find shortest path from a given vertex to all other known vertices (routers)
- The resulting set of paths forms a spanning tree
- For a graph with  $V$  vertices requires  $V-1$  iterations
- In each iteration the number of operations is proportional to  $V$
- Iterative: of  $O(V^2)$  complexity for a graph with  $V$  vertices
- Used in link state routing protocols: requires information about *topology* and *link costs*



# Dijkstra algorithm

---

Given:

- $N$ : set of vertices in graph
- $S$ : origin vertex
- $T$ : set of vertices (already) added by the algorithm
- $w(i,j)$ : cost of path from  $i$  to  $j$
- $L(n)$ : cost of the path with lower cost from  $s$  to  $n$ , already added by the algorithm

The algorithm uses 3 steps

- Step 1: initialization
- Step 2 and 3 are repeated until  $T = N$  (until spanning tree starting at origin vertex is formed )



# Dijkstra algorithm

---

## Step 1 – Initialization

1.1  $T = \{s\}$

1.2  $L(n) = w(s,n)$

$$L(s) = 0$$

$$L(i) = w(s,i) \text{ (to neighbouring nodes of } s)$$

$$L(n) = \textbf{infinite} \text{ (to other nodes which are not neighbors of } s)$$

# Dijkstra algorithm

---

## Step 2 – chooses next vertex to add

2.1 Find  $x$  not belonging to  $T$  for which  $L(x) = \min L(j)$ ,  
for all  $j$  not in  $T$

2.2 Add  $x$  to  $T$

2.3 Add link to  $x$  to  $T$

## Step 3 – update least cost paths


$L(n) = \min [ L(n), L(x) + w(x, n) ]$ ,  
for all  $n$  not in  $T$



# Dijkstra algorithm (example)

Step 1: Initialization


$T = \{V1\}$



---


$$T = \{V1, V4\}$$

### Step 3: update least cost paths




# Dijkstra algorithm (example)

Step 2: add V2 to T

$T = \{V1, V4, V2\}$

Step 3: update least cost paths




# Dijkstra algorithm (example)

Step 2: add V5 to T

$T = \{V1, V4, V2, V5\}$

Step 3: update least cost paths




# Dijkstra algorithm (example)

Step 2: add V3 to T

$T = \{V1, V4, V2, V5, V3\}$


Step 3: update least cost paths



# Dijkstra algorithm (example)

$T = \{V1, V4, V2, V5, V3, V6\} = N$


Spanning tree from node  
V1 is formed






# Distance-Vector Protocols

- Each router only knows directly-connected neighbours (reachable via the same physical link), and link costs to its neighbours
- Contrary to link-state protocols, each router possesses a limited vision of the topology



RA only "knows" RB and RC and trust those neighbours to route traffic towards other networks

Same thing for RB in relation to RD and RA




# Distance-Vector Protocols

---

- Each router sends periodic *updates* to its neighbours (even if no changes in the network have to be reported)
- The updates include the complete contents of the routing table of the router

Interface can also be  
"passive" (doesn't  
receive updates)



# Routing protocols (link-state and distance-vector)

Protocolo	Tipo de encaminhamento		Algoritmo
	interior	exterior	
<b>RIP (<i>Routing Information Protocol</i>), v1</b>	•		<i>distance-vector</i>
<b>RIP (<i>Routing Information Protocol</i>), v2</b>	•		<i>distance-vector</i>
<b>IGRP (<i>Interior Gateway Routing Protocol</i>)</b>	•		<i>distance-vector</i>
<b>EIGRP (<i>Enhanced Interior Gateway Routing Protocol</i>)</b>	•		<i>distance-vector</i> <i>+link-state</i>
<b>OSPF (<i>Open Shortest Path First</i>)</b>	•		<i>link-state</i>
<b>EGP (<i>Exterior Gateway Protocol</i>)</b>		•	—
<b>BGP (<i>Border Gateway Protocol</i>)</b>		•	—

# Router Information Protocol

---

- RIP-2 – Defined in RFC 1723 (1994)
- Router *broadcasts* routing information at each 30 s
- Sends immediate update upon detecting change on a link
- Router uses information received from its neighbours to calculate the shortest paths to all reachable (and known) destinations
- Uses *hop count* as metric
- Maximum hop count of 15 (16 is considered “infinite”, or “unreachable”)
- Reports are broadcasted to neighbours
- A route expires if no update is received for 180s

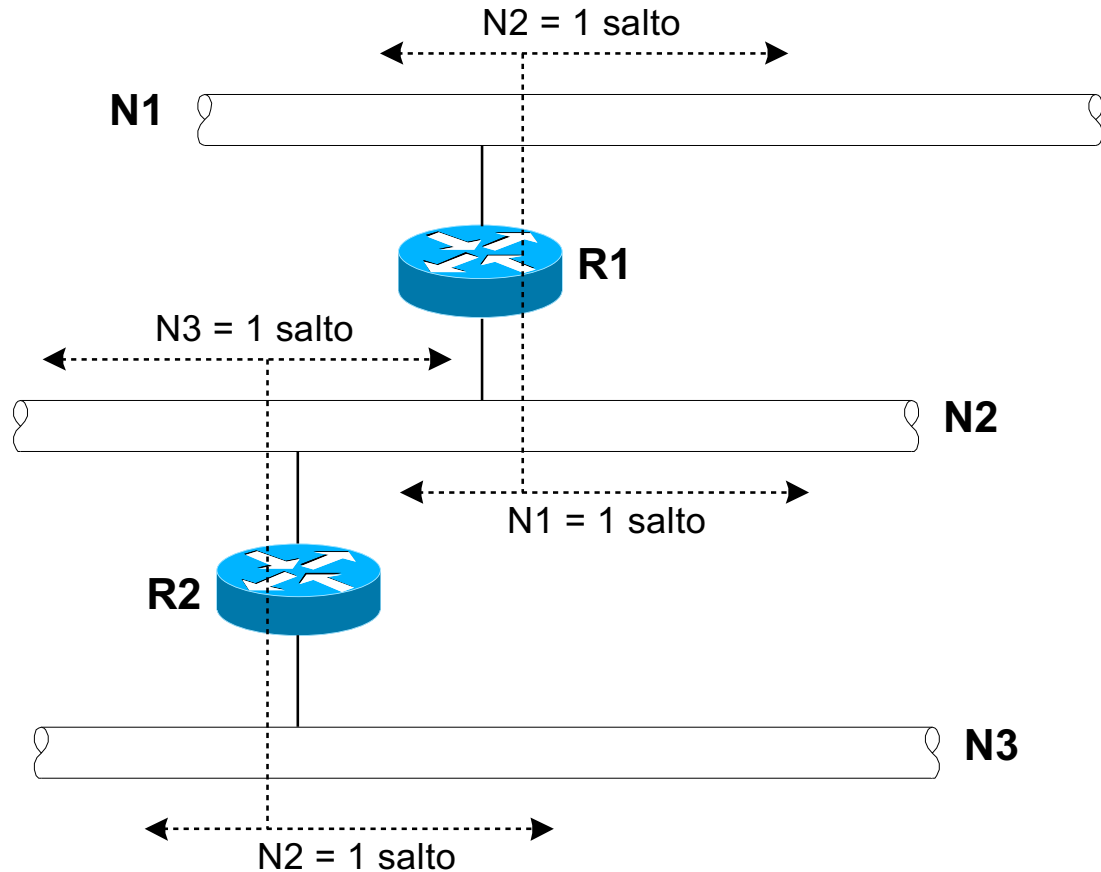
# Routing Information Protocol

---

- May store up to 6 equal cost paths to same destination
- Supports load balancing using paths with the same cost
- Timer values used by the protocol:
  - **Update interval:** 30s (periodic update of routes sent to neighbours)
  - **Invalid timer:** 180s (time since last update for route, upon which route is marked as invalid and put “on hold”: *hop count* 16 or “infinite”)
  - **Flush timer:** 240 s (time since last update upon which route is flushed or deleted)

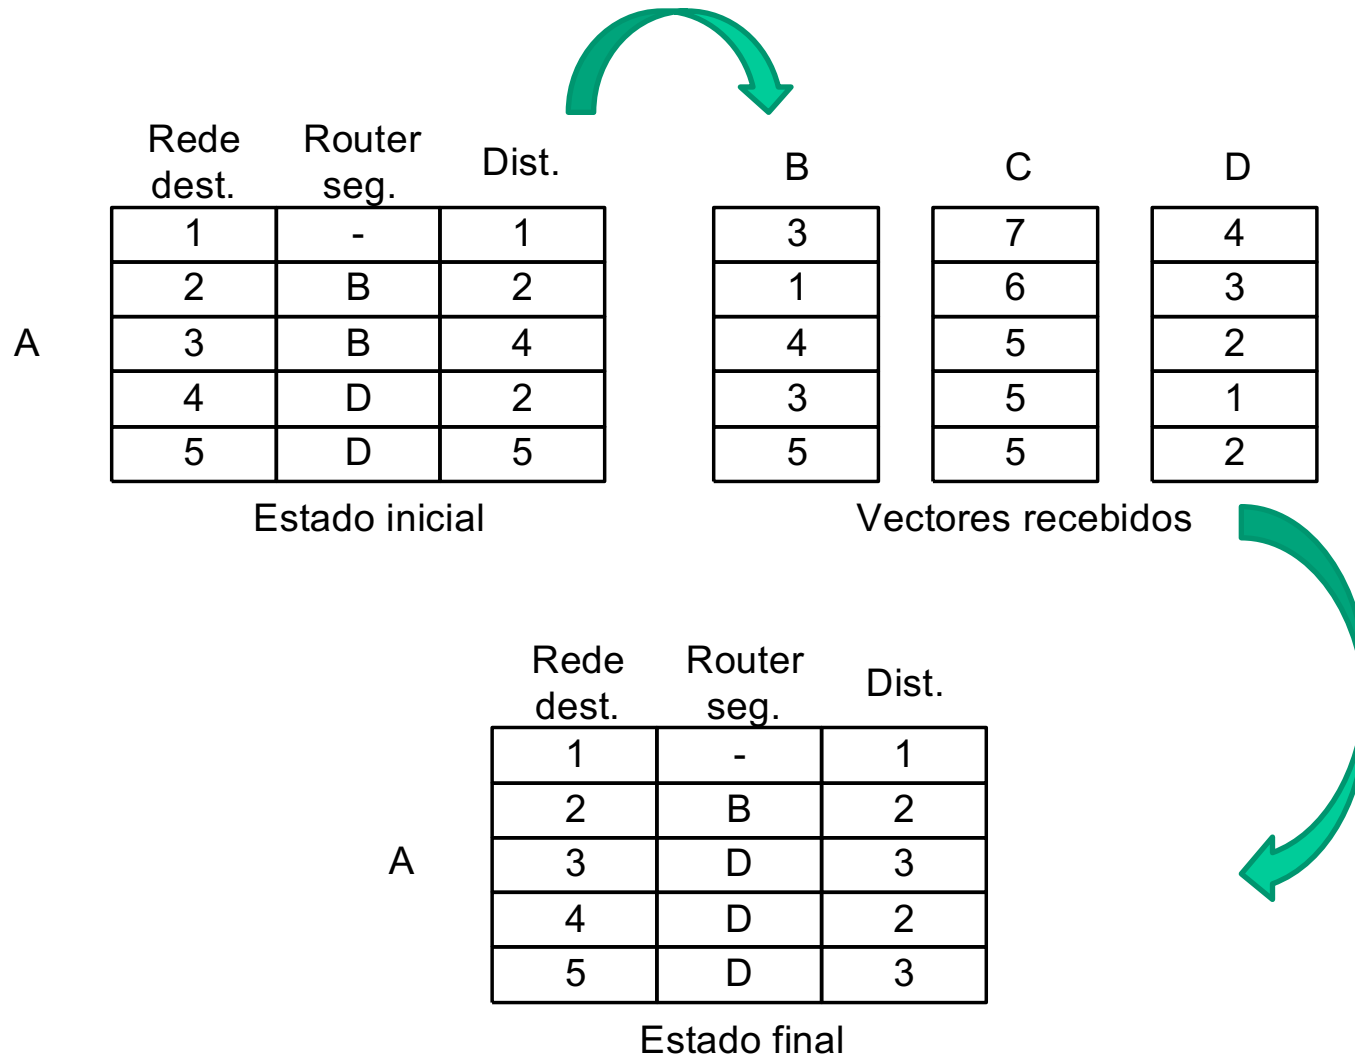
# Routing Information Protocol

---



# Routing Information Protocol

Example (update of routes using “distance vector” strategy):



# T04: outline

## 4.1 Overview of Network layer

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- service model
- datagram format
- fragmentation
- IPv4 addressing
- network address translation

## 4.4 Routing protocols

- link state
- Distance vector
- Dijkstra and RIP

## 4.5 ICMP: The Internet Control Message Protocol



# ICMP: internet control message protocol

- used by hosts & routers to communicate network-level information

- error reporting:  
unreachable host, network, port, protocol
- echo request/reply (used by ping)

- network-layer “above” IP:

- ICMP msgs carried in IP datagrams

- **ICMP message:** type, code plus first 8 bytes of IP datagram causing error

<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header


# Traceroute and ICMP

- source sends series of UDP segments to destination
  - first set has TTL = 1
  - second set has TTL=2, etc.
  - unlikely port number
- when datagram in  $n$ th set arrives to  $n$ th router:
  - router discards datagram and sends source ICMP message (type 11, code 0)
  - ICMP message include name of router & IP address

- when ICMP message arrives, source records RTTs

## *stopping criteria:*

- UDP segment eventually arrives at destination host
- destination returns ICMP “port unreachable” message (type 3, code 3)
- source stops



# T04: summary

## 4.1 Overview of Network layer

- data plane
- control plane

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- network address translation


## 4.4 Routing protocols

- link state
- Distance vector
- Internal and external routing

## 4.5 ICMP: The Internet Control Message Protocol

# T04: Bibliography

J. Kurose and K. Ross, “Computer Networking – a top-down approach”, Pearson. Chapter 4: The network Layer



# Redes de Comunicação 2023/2024

---

## T04 Network Layer Extra material

Jorge Granjal  
University of Coimbra



Slides adapted from: F Kurose and K.W. Ross, All Rights Reserved

# T04: Firewalls

- A firewall is a network security device that monitors incoming and outgoing network traffic and decides whether to allow or block specific communications.
- Firewalls support a first line of defense in network security. They establish a barrier between secured and controlled internal (trusted) networks.

## Introduction of Firewall in Computer Network

### DDoS Attack Explained

