



Multicast: Comunicação em Grupo

Sistemas Distribuídos, 2011

Multicast

- ◆ Ideia: enviar a mesma informação para diversos receptores simultaneamente, utilizando uma única mensagem


- Unicast
- Multicast
- Broadcast



S – Sender
R – Receiver


Multicast

- ◆ No entanto, pretende-se que não seja necessário utilizar uma ligação por cada receptor.
 - Nas redes locais → uso das características da Ethernet
 - Nas redes globais → uso de routers multicast




Multicast Routers

- ◆ IP packets can be multicast on a local network and on the wider Internet.
- ◆ Local multicasts make use of the hardware multicast feature of Ethernet (single message).



- ◆ Internet multicasts make use of multicast routers, which forward single datagrams to routers on other networks with members, where they are again hardware multicast to local members.



Onde é útil usar Multicasting?

- ◆ Aplicações multimédia
 - Difusão de video, som, etc.
 - Jogos multi-utilizador
- ◆ Servidores redundantes (fault-tolerance)
- ◆ Serviços de descoberta (spontaneous networking)
- ◆ Melhorar o desempenho com dados replicados
- ◆ Propagação de notificação de eventos

IP Multicast

- ◆ Routers interpret any datagram sent to an IP address in the range of:
224.0.0.0 to 239.255.255.255 as *multicast*.
- ◆ Any IP application with a UDP socket can send to a multicast address (with some limitations).
- ◆ Applications that *join* a **multicast group** can receive multicast datagrams sent to that group address.

Receiving/Sending Multicasts

- ◆ A multicast receiver application must:
 1. Get a UDP socket.
 2. Bind to the application's port number (i.e. name the socket).
 3. Join the application's multicast address group.
 4. Receive.
 5. Close the socket when complete.

- ◆ A multicast sender must:
 1. Get a UDP socket.
 2. Set the IP Time-To-Live appropriately.
 3. Send to the application's multicast address and port number.
 4. Close the socket when complete.

- ◆ The sender does not have to join a multicast group.

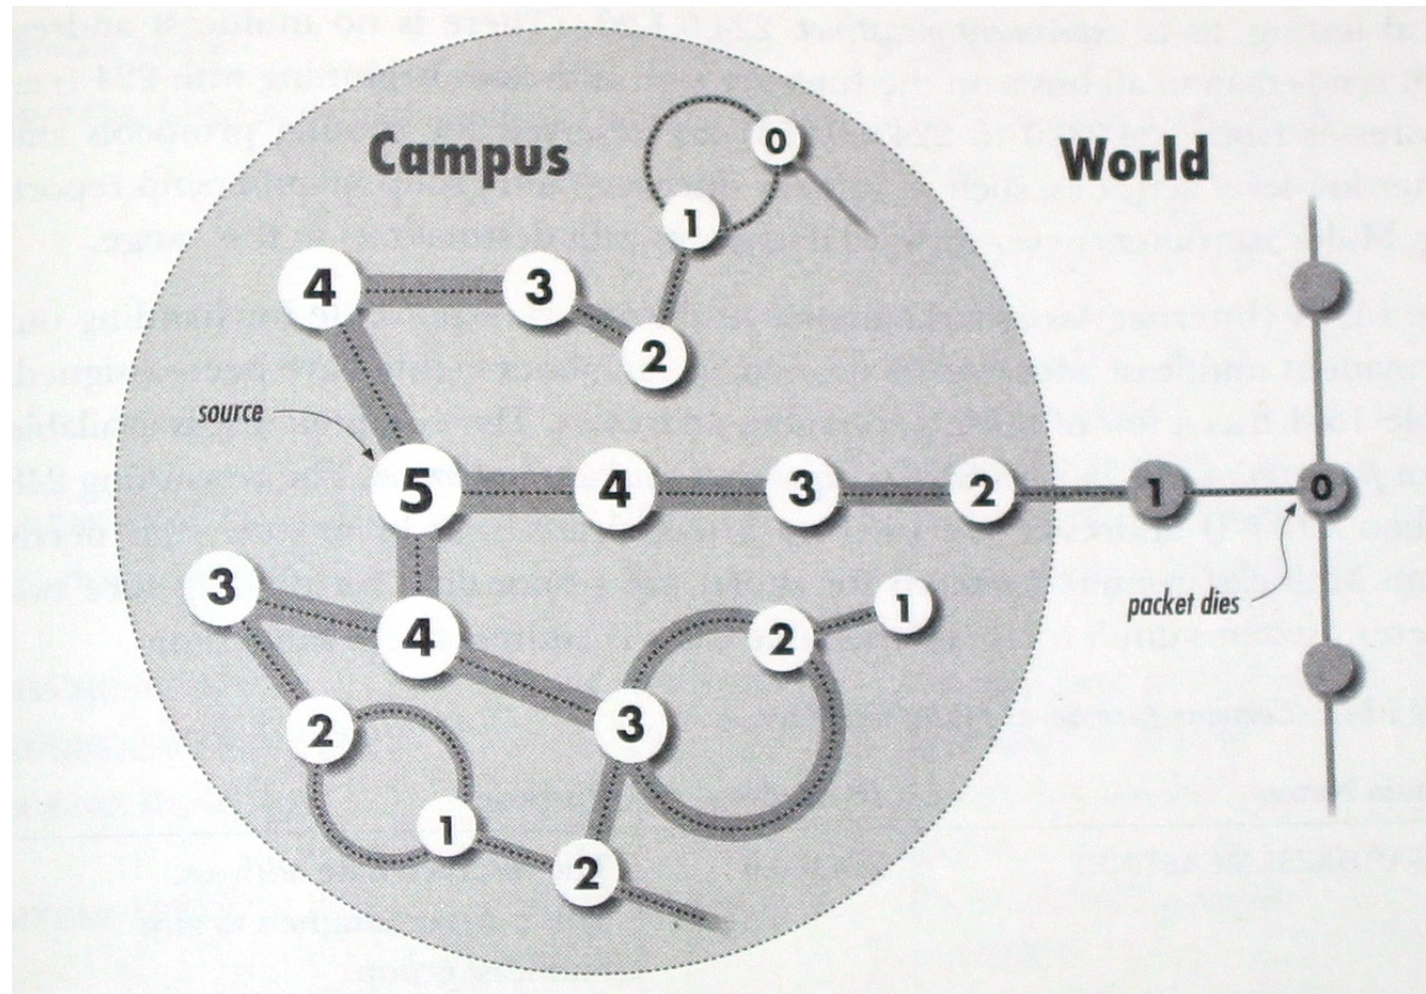
IP Multicast

- ◆ When a multicast message arrives at a computer, copies are forwarded to all of the local sockets that have joined the specific multicast address and the specified port number.
- ◆ The membership of multicast groups is dynamic: members can join and leave whenever they want.
- ◆ IP multicast is available only via UDP.
 - Unreliable: same semantics than UDP

TTL: Time-to-Live

- ◆ To limit the distance of propagation of a multicast datagram, the sender specifies the number of routers it is allowed to pass:
 - Time-to-Live: TTL
- ◆ The default value is 1.
- ◆ This allows the multicast propagate only on the local network.
- ◆ MulticastSocket
 - `public void setTimeToLive(int ttl)`



Pack dies when TTL reaches 0



Internet Multicasting

- ◆ Most providers of push media currently use ***unicast*** to deliver content to their customers. This works, but does not scale well.
- ◆ ***Multicast*** allows a content provider to send a single datastream to a single address, a datastream that network routers subsequently distribute to as many receivers as desired.
- ◆ Multicast requires no additional effort on the part of the sender to add new receivers, since the network handles the distribution from the single datastream.

The MBone!



Multicast Address Allocation

- ◆ Permanent Multicast Groups:

- 224.0.0.1 224.0.0.255

- ◆ Temporary Groups:


- 225.0.0.1 239.255.255.255

Temporary Multicast Group

- ◆ When a mcast group is created it requires a free multicast address to avoid accidental participation in an existing group.
- ◆ The IP multicast protocol does not address this issue.
- ◆ When the users only need to communicate locally they set the TTL to a small value.
- ◆ However, programs using IP multicasting over the Internet require a solution to the problem (directory of multicast sessions).


Types of multicast: One-to-Many Multicasting

- ◆ Real-time data distribution
 - Weather, stocks, telemetry, and remote sensing
- ◆ File distribution
 - Software updates, database mirrors, and web caching
- ◆ Cryptographic key distribution
- ◆ Network management
- ◆ System configuration



Types of multicast: Many-to-Many Multicasting

- ◆ Conferencing
 - Video, audio, and whiteboard sharing
- ◆ Collaborative document sharing
- ◆ Interactive distance learning
- ◆ Virtual reality
- ◆ Gaming



Java: MulticastSocket

- ◆ The multicast datagram socket class is useful for sending and receiving IP multicast packets.
- ◆ `java.net.MulticastSocket` *extends* `java.net.DatagramSocket`
 - A MulticastSocket is a (UDP) DatagramSocket, with additional capabilities for joining "groups" of other multicast hosts on the Internet.
- ◆ A multicast group is specified by a class D IP address and by a standard UDP port number.
- ◆ Class D IP addresses are in the range: 224.0.0.0 to 239.255.255.255.

Multicast peer joins a group and sends and receives datagrams

```
MulticastSocket socket = null;

try {
    InetAddress group = InetAddress.getByName("225.0.0.0");
    socket = new MulticastSocket(6789);
    socket.joinGroup(group);
    socket.setTimeToLive(1);

    byte[] buf = "Hello".getBytes();
    DatagramPacket msgOut =
        new DatagramPacket(buf, buf.length, group, port);
    socket.send(msgOut);

    ////

    byte[] inBuf = new byte[8*1024];
    DatagramPacket msgIn =
        new DatagramPacket(inBuf, inBuf.length);
    socket.receive(msgIn);

    System.out.println("Received:" + new String(inBuf, 0,
                                                msgIn.getLength()));

    socket.leaveGroup(group);
} catch (Exception e) { e.printStackTrace(); }

...
```

Example 2: Receiver

```
import java.net.*;
int port = 5000;
group = "225.4.5.6";


// Create the socket and bind it to port 'port'.
MulticastSocket s = new MulticastSocket(port);

// join the multicast group
s.joinGroup(InetAddress.getByAddress(group));

// Now the socket is set up and we are ready to receive packets
byte buf[] = new byte[1024];
DatagramPacket pack = new DatagramPacket(buf, buf.length);
s.receive(pack);

// Do something useful with the data we just received,
System.out.println("Received data from: " + pack.getAddress().toString() + ":" + pack.getPort() + "
    with length: " + pack.getLength()); System.out.write(pack.getData(),0,pack.getLength());

// And when we have finished receiving data leave the multicast group and close the socket
s.leaveGroup(InetAddress.getByAddress(group));
s.close();
```



When the socket is created the [`DatagramSocket.setReuseAddress\(boolean\)`](#) method is called to enable the `SO_REUSEADDR` socket option.

Example 2: Simple Sender (does not join group)

```
import java.net.*;

int port = 5000;
String group = "225.4.5.6";
int ttl = 1;

// Create the socket but we don't bind it as we are only going to send data
MulticastSocket s = new MulticastSocket();
// We don't have to join the multicast group if we are only sending data

// Fill the buffer with some data
byte buf[] = new byte[10];
for (int i=0; i<buf.length; i++)
    buf[i] = (byte)i;
// Create a DatagramPacket
DatagramPacket pack = new DatagramPacket(buf, buf.length,
                                         InetAddress.getByName(group), port);
// Do a send. Note that send takes a byte for the ttl and not an int.
s.send(pack,(byte)ttl);
// And when we have finished sending data close the socket
s.close();
```

Important methods of MulticastSocket

```
// Create a multicast socket
```

```
public MulticastSocket();
```

```
// Create a multicast socket in a specific port
```

```
public MulticastSocket(int port);
```

- // Get and set the current time-to-live

```
public int getTimeToLive();
```

```
public void setTimeToLive(int ttl);
```

- // Join and leave a multicast group

```
public void joinGroup(InetAddress group);
```

```
public void leaveGroup(InetAddress group);
```

Joining/Leaving Mcast Groups

- ◆ When one sends a message to a multicast group, all subscribing recipients to that host and port receive the message (within the time-to-live range of the packet).
- ◆ The socket needn't be a member of the multicast group to send messages to it.
- ◆ A socket relinquishes membership in a group by the `leaveGroup(InetAddress addr)` method.
- ◆ Multiple MulticastSocket's may subscribe to a multicast group and port concurrently, and they will all receive group datagrams.


Failure Model

- ◆ Datagram multicasts have the same failure characteristics as UDP datagrams:
 - Messages can be lost
 - Messages can arrive out-of-order
 - Messages can be duplicated
- ◆ Thereby, with IP mcast we have the following semantics:
 - Unreliable Multicast
- ◆ There are other (more powerful) multicast semantics...
 - Check [Coulouris2001] section 11.4

Reliability Semantics of Multicast


- ◆ 0-reliable
- ◆ 1-reliable
- ◆ m-out-n reliable
- ◆ All-reliable

How to achieve 1-Reliable?




Example: read data from replicated servers

How to achieve m-out-n Reliable?



Example: voting system

How to achieve All-Reliable?



Better Algorithm?...

Example: updating replicated data...

...To be continued...

- ◆ Mais à frente será leccionado um tópico sobre comunicação em grupo, onde esta questão da fiabilidade será explicada.